

Project Autonomous and Adaptive Systems 2021-22

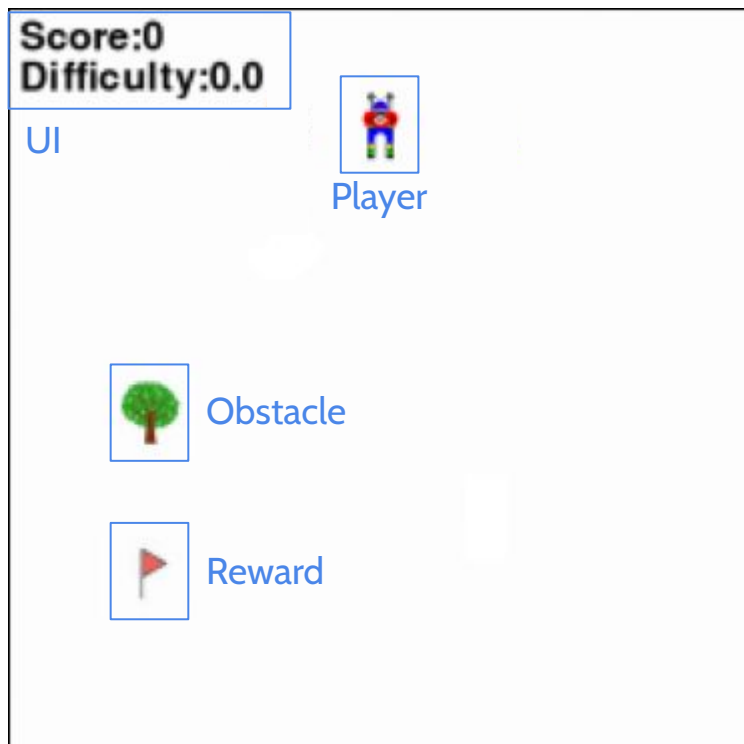
giovanni.minelli2@studio.unibo.it

Giovanni Minelli

Project objective: create a game environment parameterised by a difficulty value to enable the training of better autonomous player agents



Project objective: create a game environment parameterised by a difficulty value to enable the training of better autonomous player agents



- Interface with game score and difficulty
- Vertical scrolling environment
- 5 degrees of action for the player
- -20 score for touching a tree
- +10 score for touching a flag
- Maps generated on the fly


Two agents:

Solver: predict the player's action at every step


Generator: generates a new map of obstacles on demand

Env observation

View resized, converted to grayscale,
frame skipping 8 (max 6px), frame stacking 4



CNN: 224x224x4



MLP: 64*64*3

Actor Generator

Layer (type)	Kernel	Output Shape	Param #
input		224x224x1	0
input_aux		1x1	0
input_positions		1x100	0
conv2d	8x8 s 4	55x55x32	2080
conv2d_1	4x4 s 2	26x26x64	32832
conv2d_2	3x3 s 1	24x24x64	36928
flatten+concat		36864+1	0
dense		1x512	18875392
prediction		1x100	51300
filter_pred (Mult)		1x100	0

Actor Solver

Layer (type)	Kernel	Output Shape	Param #
input		224x224x4	0
conv2d	8x8 s 4	55x55x32	8224
conv2d_1	4x4 s 2	26x26x64	32832
conv2d_2	3x3 s 1	24x24x64	36928
flatten		36864	0
dense		1x512	18874880
prediction		1x5	2565

Networks

4 CNN networks for the estimation
of a policy and a value function
(solver and generator agent)


PPO Loss

$$L_t^{CLIP+VF+S}(\theta) = \hat{\mathbb{E}}_t[L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 S[\pi_\theta](s_t)], \quad (9)$$

where c_1, c_2 are coefficients, and S denotes an entropy bonus, and L_t^{VF} is a squared-error loss $(A_t - V_\theta(s_t))^2$

[Proximal Policy Optimization Algorithms](#)

Results



30 game episodes of baseline and expert with generator

Map	Difficulty	Game score		Total objects hit	
		Baseline	Expert	Baseline	Expert
Random map		697.5	1032.5	2792	2693
Generator agent	1 (easy)	-37.5	642.5	2718	2293
	0 (medium)	-87.5	528.3	2662	2196
	-1 (hard)	-290.8	554.2	2862	2294



Play Demo

Thank you

A2C PPO framework

PPO is a family of first-order methods (unlike TRPO) that use a few tricks to keep new policies close to old.

- PPO is a policy gradient method on-policy belongs to the family of Actor-Critic algorithms
- PPO can be used for environments with either discrete or continuous action spaces.
- and as well with environments episodic or sequential dependently by the advantage function

Loss in three terms:

$$L_t^{CLIP+VF+S}(\theta) = \hat{\mathbb{E}}_t[L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 S[\pi_\theta](s_t)], \quad (9)$$

where c_1, c_2 are coefficients, and S denotes an entropy bonus, and L_t^{VF} is a squared-error loss $(A_t - V_\theta(s_t))^2$

[Proximal Policy Optimization Algorithms](#)

$$L^{CLIP}(\theta) = \min(r_t(\theta)A_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)A_t)$$

Advantage (note it's almost the same error of critic network)

$$A_t = r_{t+1} + V(s_{t+1}) - V(s_t)$$

Ratio for policy

$$r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)}$$